

ML in the Valley

*ML insight and lessons
learned from industry
practice*

[« Back to blog](#)

Cheng-Tao Chu

Director of Analytics at
Codecademy. Specialties: data
engineering and machine
learning. Formerly: Google,
LinkedIn and Square.

[LinkedIn](#)
[Twitter](#)
[Google](#)
[Github](#)



Posted [6 years ago](#)
June 3, 2014 at 2:48 AM

65635 views

Machine Learning Done Wrong

Statistical modeling is a lot like
engineering.

In engineering, there are various ways to build a key-value storage, and each design makes a different set of assumptions about the usage pattern. In statistical modeling, there are various algorithms to build a classifier, and each algorithm makes a different set of assumptions about the data.

When dealing with small amounts of data, it's reasonable to try as many algorithms as possible and to pick the best one since the cost of experimentation is low. But as we hit "big data", it pays off to analyze the data upfront and then design the modeling pipeline (pre-processing, modeling, optimization algorithm, evaluation, productionization) accordingly.

As pointed out in my previous [post](#), there are dozens of ways to solve a given modeling problem. Each model assumes something different, and it's not obvious

how to navigate and identify which assumptions are reasonable. In industry, most practitioners pick the modeling algorithm they are most familiar with rather than pick the one which best suits the data. In this post, I would like to share some common mistakes (the don't-s). I'll save some of the best practices (the do-s) in a future post.

1. Take default loss function for granted

Many practitioners train and pick the best model using the default loss function (e.g., squared error). In practice, off-the-shelf loss function rarely aligns with the business objective. Take fraud detection as an example. When trying to detect fraudulent transactions, the business objective is to minimize the fraud loss. The off-the-shelf loss function of binary classifiers weighs false positives and false negatives equally. To align with the business objective, the loss function should not only penalize false negatives more than false positives, but also penalize each false negative in proportion to the dollar amount. Also, data sets in fraud detection usually contain highly imbalanced labels. In these cases, bias the loss function in favor of the rare case (e.g., through up/down sampling).

2. Use plain linear models for non-linear interaction

When building a binary classifier, many practitioners immediately jump to logistic regression because it's simple. But, many also forget that logistic regression is a linear model and the non-linear interaction

among predictors need to be encoded manually. Returning to fraud detection, high order interaction features like "billing address = shipping address and transaction amount < \$50" are required for good model performance. So one should prefer non-linear models like SVM with kernel or tree based classifiers that bake in higher-order interaction features.

3. Forget about outliers

Outliers are interesting. Depending on the context, they either deserve special attention or should be completely ignored. Take the example of revenue forecasting. If unusual spikes of revenue are observed, it's probably a good idea to pay extra attention to them and figure out what caused the spike. But if the outliers are due to mechanical error, measurement error or anything else that's not generalizable, it's a good idea to filter out these outliers before feeding the data to the modeling algorithm.

Some models are more sensitive to outliers than others. For instance, AdaBoost might treat those outliers as "hard" cases and put tremendous weights on outliers while decision tree might simply count each outlier as one false classification. If the data set contains a fair amount of outliers, it's important to either use modeling algorithm robust against outliers or filter the outliers out.

4. Use high variance model when $n \ll p$

SVM is one of the most popular off-the-shelf modeling algorithms and one of its most powerful features is the ability to fit the model with different kernels. SVM kernels can be thought of as a way to automatically combine existing features to form a richer feature space. Since this power feature comes almost for free, most practitioners by default use kernel when training a SVM model. However, when the data has $n \ll p$ (number of samples \ll number of features) -- common in industries like medical data -- the richer feature space implies a much higher risk to overfit the data. In fact, high variance models should be avoided entirely when $n \ll p$.

5. L1/L2/... regularization without standardization

Applying L1 or L2 to penalize large coefficients is a common way to regularize linear or logistic regression. However, many practitioners are not aware of the importance of standardizing features before applying those regularization.

Returning to fraud detection, imagine a linear regression model with a transaction amount feature. Without regularization, if the unit of transaction amount is in dollars, the fitted coefficient is going to be around 100 times larger than the fitted coefficient if the unit were in cents. With regularization, as the L1 / L2 penalize larger coefficient more, the transaction amount will get penalized more if the unit is in dollars. Hence, the regularization is biased and tend to penalize features in smaller scales. To mitigate the problem,

standardize all the features and put them on equal footing as a preprocessing step.

6. Use linear model without considering multi-collinear predictors

Imagine building a linear model with two variables X_1 and X_2 and suppose the ground truth model is $Y = X_1 + X_2$. Ideally, if the data is observed with small amount of noise, the linear regression solution would recover the ground truth. However, if X_1 and X_2 are collinear, to most of the optimization algorithms' concerns, $Y = 2 * X_1$, $Y = 3 * X_1 - X_2$ or $Y = 100 * X_1 - 99 * X_2$ are all as good. The problem might not be detrimental as it doesn't bias the estimation. However, it does make the problem ill-conditioned and make the coefficient weight uninterpretable.

7. Interpreting absolute value of coefficients from linear or logistic regression as feature importance

Because many off-the-shelf linear regressor returns p-value for each coefficient, many practitioners believe that for linear models, the bigger the absolute value of the coefficient, the more important the corresponding feature is. This is rarely true as (a) changing the scale of the variable changes the absolute value of the coefficient (b) if features are multi-collinear, coefficients can shift from one feature to others. Also, the more features the data set has, the more likely the features are multi-collinear and the less reliable to interpret the feature importance by coefficients.

So there you go: 7 common mistakes when doing ML in practice. This list is not meant to be exhaustive but merely to provoke the reader to consider modeling assumptions that may not be applicable to the data at hand. To achieve the best model performance, it is important to pick the modeling algorithm that makes the most fitting assumptions -- not just the one you're most familiar with.

If you like the post, you can follow me ([@chengtao_chu](#)) on [Twitter](#) or subscribe to my blog "[ML in the Valley](#)". Also, special thanks Ian Wong ([@ihat](#)) for reading a draft of this.

Upvote 62

[Tweet](#)

Like this post? [Subscribe by email »](#)

86 responses

$Y=100 \cdot X_1 - 99 \cdot X_2$

98

— charles

nice post!

as for #6, #7, would you please refer to some literature? and in particular for #7 a), would it still be the case if standardization is applied to the features?

— jyangtum

@jyangtum, I actually don't really have a good literature reference other than wikipedia in this case. For #7, standardizing features could definitely help. However, you are making the assumption of the feature is more or less normally distributed. Other

than that, you should still be careful with the colinearity among other features.

— Cheng-Tao Chu

one follow up question on the colinearity. Would you generalize the statement to log-linear model for structure prediction? e.g. CRFs

— jyangtum

Hi,

Most of your points are really very valid. However for points 2 and 4.

2: One can use non linear interactions in linear model too. Basis functions and kernel does exactly that. To build a polynomial regression. regression splines (which use the basis function) is the way to model non-linear interaction of features with target. One can do the same for logistic too and can make it non-linear. One doesn't really have to go to SVM to do that unless you are getting a better performance with SVM

4). Again, if n

— Manish Tripathi

Nice list. Even when dealing with a small volume of data, it is worth thinking about these pitfalls beforehand.

— Mathieu

Very interesting post. I would actually like to see more blogs and even research papers showing common mistakes and negative results as I believe this is the best way of learning.

I would also add incorrect sampling and wrong assumptions based on the data to this list of common mistakes.

In addition, I would suggest this video about common problems in machine learning from one of the data scientist from Kaggle: Machine Learning Gremlins - <https://www.youtube.com/watch?v=tleeC-KIsKA&fea...>

— Miguel Martinez-Alvarez

@jyangtum, unfortunately, i'm not familiar with CRFs enough to comment

— Cheng-Tao Chu

@Manish yes, and one thing worth mentioning is that even if you can do whatever feature engineering or apply whatever kernel upfront, one key benefit of using SVM with kernel is that the speed of the optimization is fast.

— Cheng-Tao Chu

nice post and Learned a lot

— Hsuan-chun

Nice post.

Would you mind if I translate this post into Korean and upload at my blog?

— Hye Jin

Hye Jin - Sure, as long as you cite the link back to the origin post

somewhere at the top of the page :) Also, do you mind sharing the link to me after you published your translated post?

.

— Cheng-Tao Chu

I have some problems with the 4th section (n

— Marc Claesen

Marc Claesen- great comment and I should have made it more clear in the post. What I tried to express was more of not using kernel rather than not using SVM when n

— Cheng-Tao Chu

Great post. Would like to see more of these and ones that compare different methods in detail. Many times the assumptions are implicit and lot of literature does not refer to them. Its only when someone has deep knowledge of the approaches and has used them in the real world that they are able to articulate these. And even then, often its based on empirical evidence, rather than with sound theoretical underpinnings. Looking forward to more posts like this!

— Anjali

A posthaven user upvoted this post.

Logistic regression is not a linear model. And where did you say you went to school? And what have you discovered by way of actionable information extracted from data in your various business postings? I would like to know because someone who thinks logistic regression is linear is starting off on a very wrong foot.

— Bill Luker Jr, PhD

Bill Luke Jr, PhD, you might have a different definition or interpretation of what a linear model is. What I described in the post use similar concept of Generalized Linear Model and more details can be found at

https://en.wikipedia.org/wiki/Linear_regression
[https://en.wikipedia.org/wiki/Generalized_linear...](https://en.wikipedia.org/wiki/Generalized_linear_model)

— Cheng-Tao Chu

Great post. shouldn't 4) be "avoid high variance" instead of "use high variance"?

— Fazlan

@Fazlan, you are right and what i was trying to say is that practitioners use high variance model when n

— Cheng-Tao Chu

Hi, I'd like to translate this great essay to Chinese. Can you give me the permission? The translated text will be published on ai.jqr.com and related Chinese social network accounts. Thanks.

— weakish

@weakish, sure as long as you cite the website and my name

— Cheng-Tao Chu

@chengtao_chu Just finished Chinese translation: <https://www.jqr.com/article/000372> Your name is cited at the beginning (in translator's note) and there is a backlink at the end of the translated text. If you have other requirements, please let me know. :-)

— weakish

Thanks for your post.

Even though this post was written 5 years ago, it seems that still so many people mindlessly use ML without second thought even at the moment. Insightful, concise, organised!

Lastly, would it be possible to translate your well-written post into Korea.

Of course, I'll drop you my link here.

— JIN

@jin, sure as long as you cite the website and my name. Please also feel free to post the link to your page in the comment once it's done. Thank you.

— Cheng-Tao Chu

61 visitors upvoted this post.

Your Name

Email

[Add Website URL »](#)

Your Comment

☐ Notify me by email when new comments are added

Comment